Why `@StateHint` POJO with `Map` or `List` Are Sensitive to "Extremely Large State"

> **TL;DR:**
> A `Map` inside `@StateHint` is not a map—it's a **serialized blob**.
> Every update rewrites the entire thing → performance collapse → 💥 fails at ~2GB **serialized value size**.

The issue isn't just "large state."

It's **how that state is physically stored and accessed**.

Table of Contents

- 1.0 What `@StateHint` Really Means Under the Hood
- 2.0 Why This Can Become a Scaling Problem
  - 2.1 Hard limit cliff at 2 GB per serialized value
- 3.0 The Solution: PTF should use `MapView` and `ListView`
  - 3.1 Why Confluent's PTF Early Access doesn't support `MapView` or `ListView` yet
- 4.0 Summary

1.0 What `@StateHint` Really Means Under the Hood

```
flowchart LR

subgraph BAD["❌ ValueState (Blob Storage)"]
    A1["Partition Key"]
    A2["Serialized POJO (Map/List)"]
    A1 --> A2

    note1["Every update:
    - Read entire blob
    - Deserialize
    - Modify
    - Serialize entire blob
    - Write back"]
end

subgraph GOOD["✅ MapState (Per-Entry Storage)"]
    B1["Partition Key + key1"] --> V1["value1"]
    B2["Partition Key + key2"] --> V2["value2"]
    B3["Partition Key + key3"] --> V3["value3"]

    note2["Every update:
    - Access single entry
    - No full deserialization
    - Incremental write"]
end
```

When you annotate a POJO with `@StateHint`, Flink backs it with a single **`ValueState<YourPojo>`**. That means the entire POJO — including any `Map` or `List` fields inside it — is treated as one atomic value.

So, this:

```
public static class MyState {
    public Map<String, Integer> myMap;
}
```

...is effectively:

```
ValueState<Map<String, Integer>> myState;
```

Implication

Every access becomes:

1. Read entire POJO from RocksDB (deserialize)
2. Modify one tiny part
3. Write entire POJO back (serialize)

Even if you only update one key in the map, you still pay the cost of the entire structure.

2.0 Why This Can Become a Scaling Problem

States that use `ListState`, a merge operation in RocksDB, can silently accumulate value sizes > 2^31 bytes and will then fail on their next retrieval.

You see in RocksDB there are two main write patterns for state:

1. Put - overwrites the entire value for a key. This is what happens with `ValueState` and `MapState` use. Each write replaces the previous value completely.

2. Merge - appends data to an exisiting key without reading-then-writing. This is what `ListState` uses. When you call `listState.add(newElement)`, Flink doesn't read the exisitng list, deserialize it, append, re-serialize, and write back. Instead, it uses RocksDB's merge operator to just append the new element directly in RocksDB. This is much more efficient for large lists.

But there is a problem with the `Merge` pattern when your state values get very large. If you have a `ListState` that grows so big that its serialized form exceeds 2^31 bytes=2GB (the max size of a Java `byte[]`), then you get silent corruption. The write appears to succeed, but the next time you read it, it crashes because the data is corrupted.

2.1 Hard limit cliff at 2 GB per serialized value

RocksDB's JNI (Java Native Interface) bridge — the layer that lets Java talk to RocksDB's native C++ code — passes data back and forth as `byte[]` arrays. Java's `byte[]` has a maximum length of `2^31 - 1` bytes (just under 2GB).

So when Flink tries to serialize your entire POJO (including its map or list) into a `byte[]` to hand off to RocksDB, if the serialized result exceeds 2GB, it silently overflows. You don't get an error immediately — the write may appear to succeed. The crash happens on the **next read**, when Flink tries to deserialize the corrupted bytes back into your POJO.

That's what makes it particularly nasty — it's not a clean "your state is too big" error at write time. It's a silent corruption that blows up later, potentially after a checkpoint/restore cycle, making it hard to diagnose.

`MapState` is used as a replacement for `ListState` or `ValueState` in case the records get too big for the RocksDB JNI bridge.

3.0 The Solution: PTF should use `MapView` and `ListView`

`MapView` and `ListView` are facades over Flink's native **`MapState`** and **`ListState`** primitives. In RocksDB:

- Each **`MapState` entry** is stored as an independent RocksDB key (`partition_key + map_entry_key → value`). You can look up, update, or delete a single entry without touching the rest.
- Each **`ListState` entry** is similarly stored per-element.

This is why `MapView` and `ListView` are designed for "extremely large" collections — you never materialize the whole thing on the JVM heap. You do surgical point lookups via JNI into RocksDB.

3.1 Why Confluent's PTF Early Access doesn't support `MapView` or `ListView` yet

This is a tracked sub-task: FLINK-37598 — "Support `ListView` and `MapView` in PTFs" — filed by Timo Walther (the FLIP-440 author) to add list state and map state support in PTFs. It wasn't in the original FLIP-440 scope and is a planned addition. The Confluent Early Access simply reflects the upstream Flink state: `MapView` and `ListView` in PTFs are not yet implemented in Flink itself, let alone wired through Confluent's managed platform.

4.0 Summary

When you use `@StateHint` with a POJO containing a `Map` or `List` field, Flink treats the **entire POJO as one single value** in storage. Every time an event arrives, Flink has to:

1. Read the whole thing from RocksDB and deserialize it into memory
2. Let your code make whatever change it needs (maybe just updating one entry)
3. Serialize the whole thing back and write it to RocksDB

So if your map has 10,000 entries and you only need to update one of them, you're still reading and writing all 10,000 entries on every single event.

`MapView` and `ListView` (which aren't supported in PTFs yet) would fix this by storing each map and list entry as its own independent record in RocksDB — so you only touch the one entry you actually need.

The phrase "extremely large state" in the Confluent docs is basically shorthand for: *"at some point your map or list gets big enough that this full serdes cycle on every event becomes a real performance problem"* — and there's also a hard technical cliff at 2GB where the whole thing crashes.